



Fall 2021, Computer Vision
Homework 1

by 64160010 - Rumeysa ÇELİK

November 4, 2021

Problem Set 1

Problem 1: (Basic image operations)



Figure 1: Input Image

A Python script was written to crop a region from the center; region size set to be half the size of the input image. Image saved as png file. The red channel of the image was extracted and displayed. The image has been converted to grayscale and displayed. Sobel filters in the x and y directions were defined in the code. These filters were applied to the grayscale image and the results were displayed. Gradient magnitude and gradient direction were obtained using these gradients. Found a way to show gradient direction; The gradient magnitude and gradient direction are displayed. Laplacian of Gaussian images were obtained for different sigma values and the results were displayed.

Code:

```
1 # -*- coding: utf-8 -*-
2 """
3 Created on Fri Oct 31 19:11:12 2021
4
5 @author: Rumeysa CELIK
6 """
7 from PIL import Image
8 import cv2
9 import matplotlib.pyplot as plt
10 import numpy as np
11
12
13 #PART A: Write a Python script to crop out a region from the center; the
14 #          region size      should be half the size of the input image. Save
15 #          the image as a pngfile.
16 img= cv2.imread("/Users/rumeysacelik/Desktop/ComputerVisionHW1/huser.
jpeg",cv2.COLOR_BGR2GRAY)
```

```

17 cv2.imshow('original image',img)
18
19 print("ORIGINAL IMAGE SHAPE:",img.shape) # Print image shape
20 #Cropping an image
21
22 width, height = img.shape[1], img.shape[0]
23 c_x,c_y=int(width/2),int(height/2)
24 mid_x, mid_y = int(width/2), int(height/2)
25
26
27 cropped_img = img[c_y-int(c_y/2):c_y+int(c_y/2),c_x-int(c_x/2):c_x+int(
    c_x/2)]
28 print("CROPPED IMAGE SHAPE:",cropped_img.shape)
29
30 # Display cropped image
31 cv2.imshow("cropped image", cropped_img)
32
33 # Save the cropped image
34 cv2.imwrite("output/Cropped_Image.png", cropped_img)
35
36
37
38 #PART B: Extract the red channelof the imageand display it.
39
40 red_img= img[:, :, 2]
41 cv2.imshow("RED_image", red_img)
42 # Save the red image
43 cv2.imwrite("output/RED_image.png", red_img)
44
45
46 #PART C: Convert the image to grayscaleand display it.
47
48 gray_img= cv2.cvtColor(img, cv2.COLOR_BGR2GRAY) #converts BGR to gray
49 cv2.imshow('GRAY_img',gray_img)
50 # Save the red image
51 cv2.imwrite("output/GRAY_img.png", gray_img)
52
53 #PART D: Define the Sobel filters in x and y directionsin your code.
      Apply thesefilters to the grayscale image and displaythe results.
      Using these gradients, obtain the gradient magnitude and gradient
      orientation. Find out a way to display the gradient orientation;
      display the gradient magnitude and gradient orientation.
54
55
56 gX = cv2.Sobel(gray_img, ddepth=cv2.CV_64F, dx=1, dy=0, ksize=5)
57 gY = cv2.Sobel(gray_img, ddepth=cv2.CV_64F, dx=0, dy=1, ksize=5)
58
59 # scaling the gradients
60 gX = cv2.convertScaleAbs(gX)
61 gY = cv2.convertScaleAbs(gY)
62
63
64 # combining the gradient representations into a single image with
      addWeighted by giving them equal weight
65 combined = cv2.addWeighted(gX, 0.5, gY, 0.5, 0)

```

```

66 # show our output images
67 cv2.imshow("Sobel_X", gX)
68 cv2.imshow("Sobel_Y", gY)
69 cv2.imshow("Sobel_Combined", combined)
70 # Save the sobel images
71 cv2.imwrite("output/Sobel_X.png", gX)
72 cv2.imwrite("output/Sobel_Y.png", gY)
73 cv2.imwrite("output/Sobel_Combined.png", combined)
74
75 # computing the gradient magnitude and orientation
76 magnitude = np.sqrt((gX ** 2) + (gY ** 2))
77 orientation = np.arctan2(gY, gX) * (180 / np.pi) % 180 #converting to
    degree
78
79
80 print("MAGNITUDE\n",magnitude)
81 print("ORIENTATION\n",orientation)
82
83 #converting the arrays from float16 to uint8 to display
84 magnitude =magnitude.astype(np.uint8)
85 orientation =orientation.astype(np.uint8)
86
87 print("MAGNITUDE\n",magnitude)
88 print("ORIENTATION\n",orientation)
89 # Save the gradient magnitude and orientation
90 cv2.imwrite("output/magnitude.png", magnitude)
91 cv2.imwrite("output/orientation.png", orientation)
92
93
94 # initializing a [1,3] figure to display the input grayscale image
    along with
95 # the gradient magnitude and orientation representations, respectively
96 (fig, axs) = plt.subplots(nrows=1, ncols=3, figsize=(16, 12))
97 # plot each of the images
98 axs[0].imshow(gray_img, cmap="gray")
99 axs[1].imshow(magnitude, cmap="jet")
100 axs[2].imshow(orientation, cmap="jet")
101 # set the titles of each axes
102 axs[0].set_title("Grayscale")
103 axs[1].set_title("Gradient Magnitude")
104 axs[2].set_title("Gradient Orientation [0, 180]")
105 # loop over each of the axes and turn off the x and y ticks
106 for i in range(0, 3):
107     axs[i].get_xaxis().set_ticks([])
108     axs[i].get_yaxis().set_ticks([])
109 # show the plots
110 plt.tight_layout()
111 plt.savefig('/Users/rumeysacelik/Desktop/ComputerVisionHW1/gradient.png
    ')
112 plt.show()
113
114 #PART E: Obtain Laplacian of Gaussian imagefor different sigma values ,
    and displaythe results.
115

```

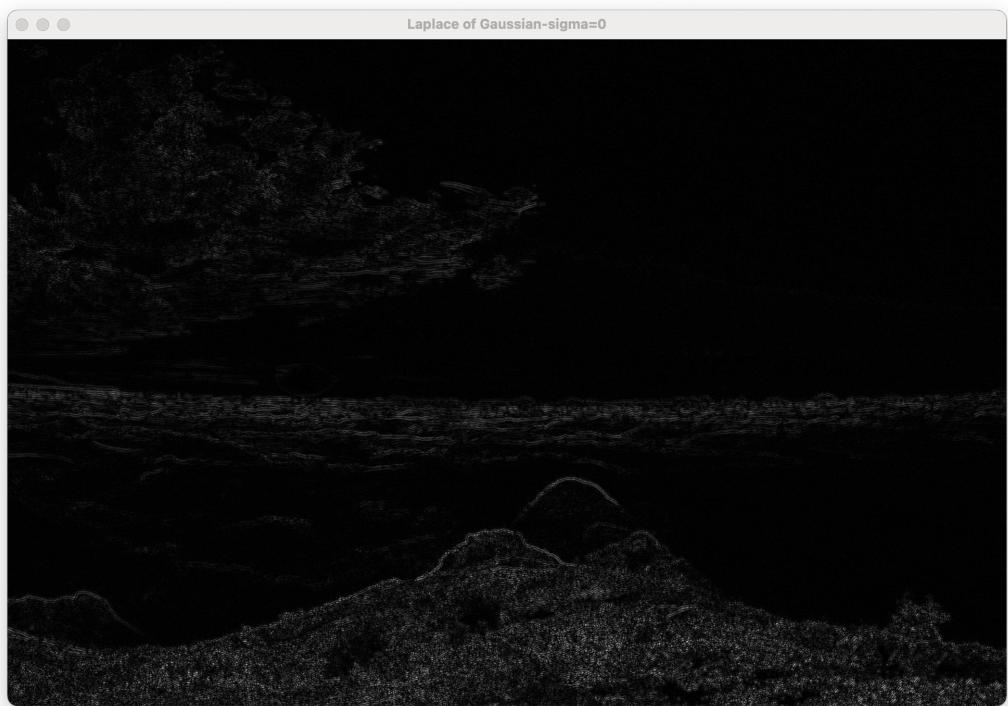
```

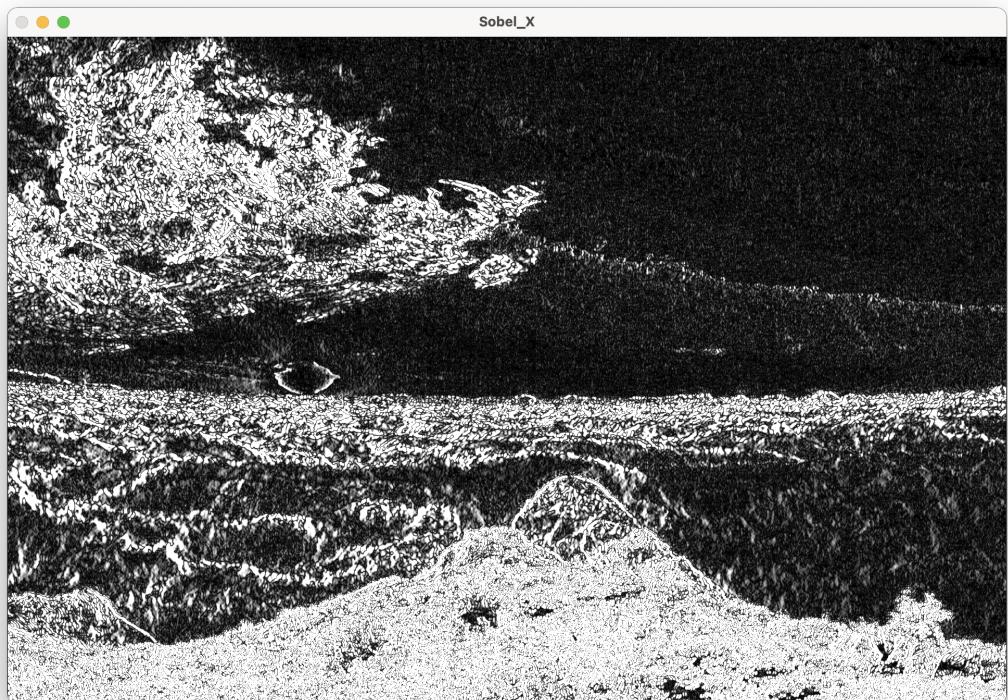
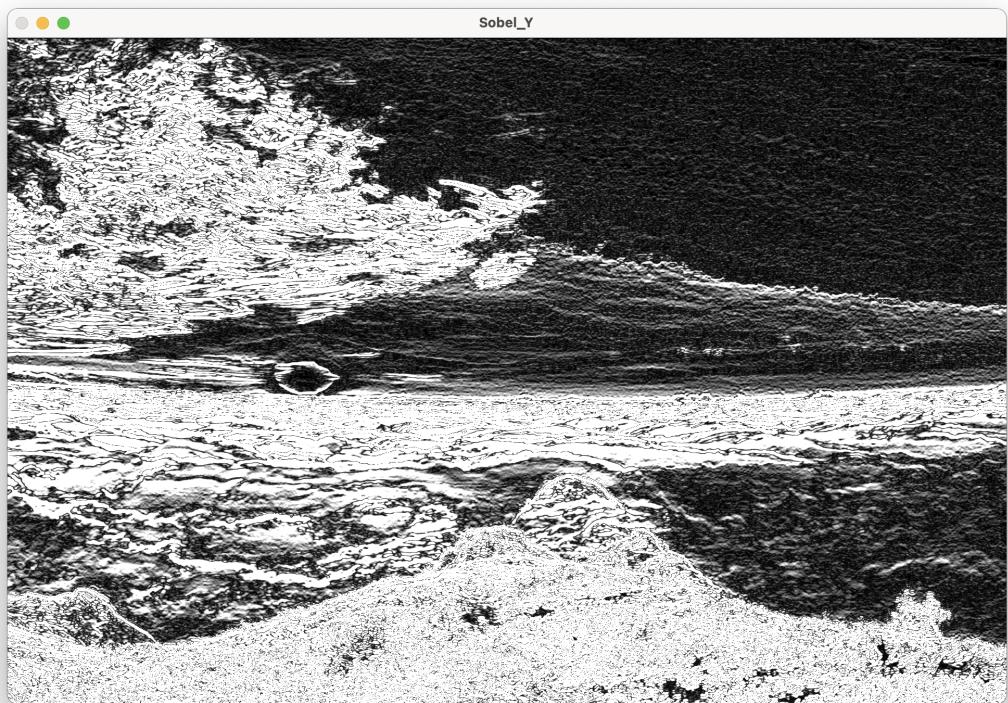
116 """Laplacian of Gaussian (LoG) is same as applying gaussian smooothing
    filter first and then
117 computing laplacian of the result"""
118 # Removing noise by blurring with a Gaussian filter
119 sigma=0#####
120 src = cv2.GaussianBlur(img, (3, 3),sigma)
121 # Converting the image to grayscale
122 src_gray = cv2.cvtColor(src, cv2.COLOR_BGR2GRAY)
123 # [laplacian]
124 # Applying Laplace function
125 dst = cv2.Laplacian(src_gray, cv2.CV_16S, ksize=3)
126 # converting back to uint8
127 log = cv2.convertScaleAbs(dst)
128 # Save the log image
129 cv2.imwrite("output/Laplacian of Gaussian(sigma=0).png", log)
130
131 cv2.imshow("Laplace of Gaussian-sigma=0", log)
132
133
134 sigma=15#####
135 # Removing noise by blurring with a Gaussian filter
136 src = cv2.GaussianBlur(img, (3, 3),sigma)
137 # Converting the image to grayscale
138 src_gray = cv2.cvtColor(src, cv2.COLOR_BGR2GRAY)
139 # [laplacian]
140 # Applying Laplace function
141 dst = cv2.Laplacian(src_gray, cv2.CV_16S, ksize=3)
142 # converting back to uint8
143 log = cv2.convertScaleAbs(dst)
144 # Save the log image
145 cv2.imwrite("output/Laplacian of Gaussian(sigma=15).png", log)
146
147 cv2.imshow("Laplace of Gaussian-sigma=15", log)
148
149 cv2.waitKey(0)
150 cv2.destroyAllWindows()

```

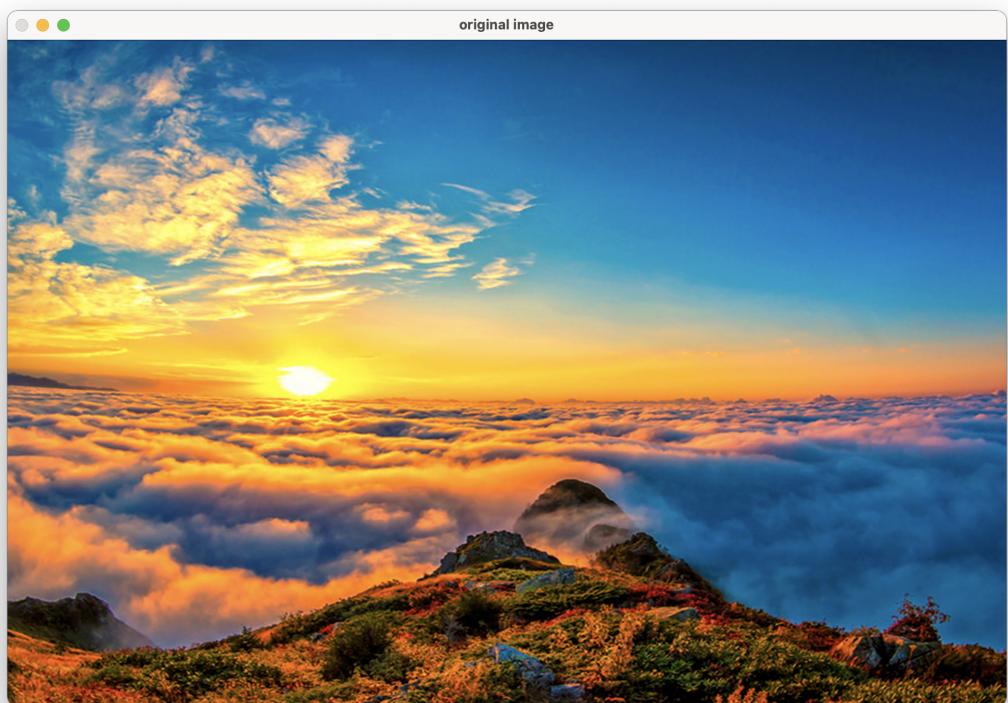
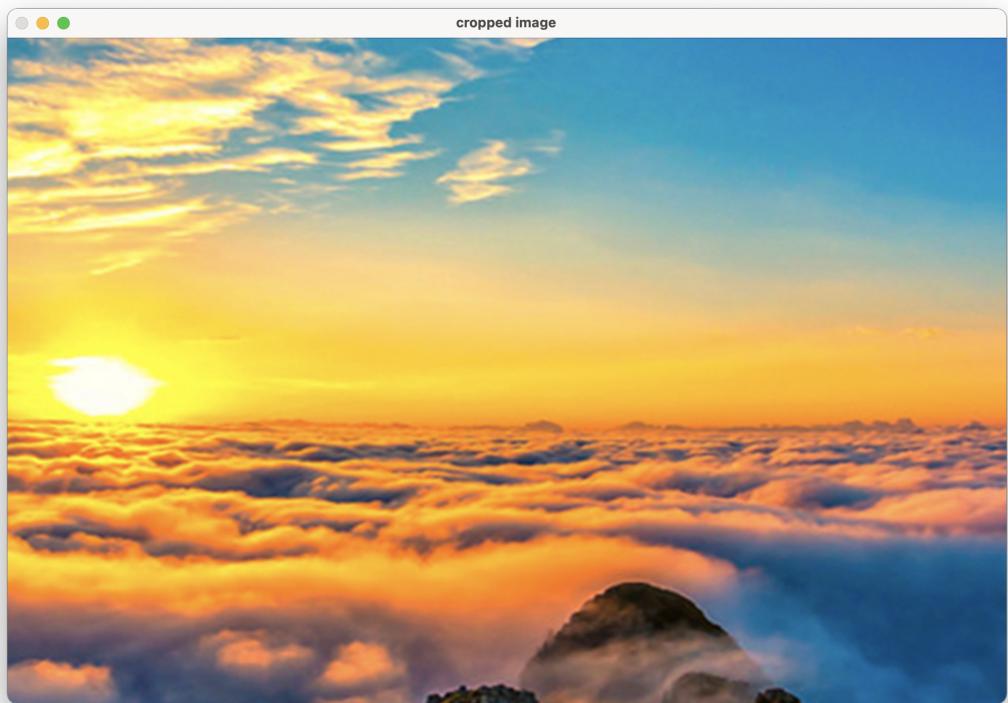
All outputs:

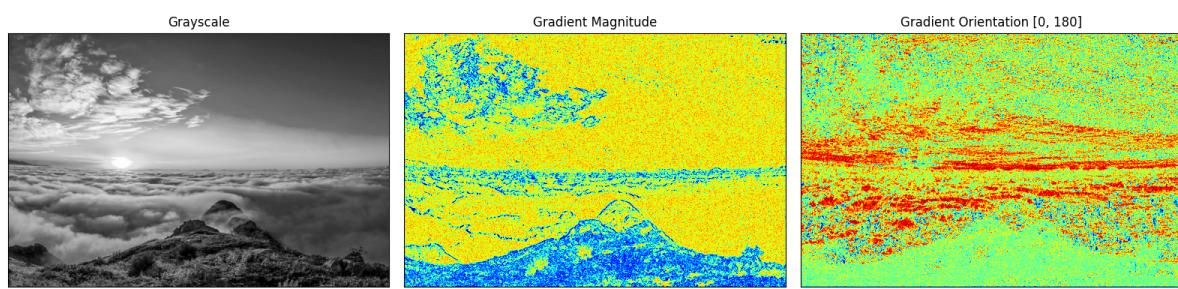












```

ComputerVisionHW1 — Python ps1_1.py — 93x44
Last login: Thu Nov  4 12:01:42 on ttys000
rumeyssacelik@Rumeysas-MacBook-Pro ComputerVisionHW1 % python ps1_1.py
Traceback (most recent call last):
  File "/Users/rumeysacelik/Desktop/ComputerVisionHW1/ps1_1.py", line 9, in <module>
    import matplotlib.pyplot as plt
ModuleNotFoundError: No module named 'matplotlib'
rumeyssacelik@Rumeysas-MacBook-Pro ComputerVisionHW1 % python ps1_1.py
[zsh: command not found: ps1_1.py
rumeyssacelik@Rumeysas-MacBook-Pro ComputerVisionHW1 % python ps1_1.py
|ORIGINAL IMAGE SHAPE: (1280, 1920, 3)
|CROPPED IMAGE SHAPE: (640, 960, 3)
MAGNITUDE
[[ 0.     4.     2.     ... 10.    12.    0.    ]
[15.1   8.25  14.42 ... 14.42  13.64  8.    ]
[ 4.    14.42  7.21 ... 14.766 8.6   6.    ]
...
[ 1.     4.125  1.     ... 8.06   1.414  1.    ]
[ 1.     8.305  14.04 ... 1.414   1.414  1.    ]
[ 0.     2.     4.     ... 1.     1.     0.    ]]
ORIENTATION
[[ 0.     0.     0.     ... 0.     0.     0.    ]
[89.94 79.4  80.5  ... 56.3  25.34  89.94]
[89.94 33.7  24.77 ... 80.8  54.47  89.94]
...
[89.94 70.1  17.4  ... 35.8  44.97  89.94]
[89.94 85.94 86.8  ... 44.97  44.97  89.94]
[ 0.     0.     0.     ... 0.     0.     0.    ]]
MAGNITUDE
[[ 0 4 2 ... 10 12 0]
[15 8 14 ... 14 13 8]
[ 4 14 7 ... 14 8 6]
...
[ 1 4 1 ... 8 1 1]
[ 1 8 14 ... 1 1 1]
[ 0 2 4 ... 1 1 0]]
ORIENTATION
[[ 0 0 0 ... 0 0 0]
[89 79 80 ... 56 25 89]
[89 33 24 ... 80 54 89]
...
[89 70 17 ... 35 44 89]
[89 85 86 ... 44 44 89]
[ 0 0 0 ... 0 0 0]]

```

Problem 2: (Basic video operations)

A Python script was written to receive a live video stream from the computer's webcam, the gradient size was calculated on the grayscale version of the image, and its input and gradient size displayed on the screen.

Code:

```

1 # -*- coding: utf-8 -*-
2 """
3 Created on Fri Oct 31 19:11:12 2021
4
5 @author: Rumeysa CELIK
6 """
7 # -*- coding: utf-8 -*-
8 """
9 Created on Sat Nov 1 01:54:05 2021
10
11 @author: Rumeysa CELIK
12 """
13 import cv2
14 import numpy as np
15 import time
16

```

```

17 #PART A: Write a Python script to take live video stream from the
18     # webcam of your computer, compute the gradient magnitude on the
19     # grayscale version of the image, and display the input and the
20     # gradient magnitude on the screen.
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66

```

#PART A: Write a Python script to take live video stream from the
 # webcam of your computer, compute the gradient magnitude on the
 # grayscale version of the image, and display the input and the
 # gradient magnitude on the screen.

```

# define a video capture object
vid = cv2.VideoCapture(0)
# used to record the time when we processed last frame
prev_frame_time = 0

# used to record the time at which we processed current frame
new_frame_time = 0

# font which we will be using to display FPS
font = cv2.FONT_HERSHEY_SIMPLEX

while True:

    # Capture the video frame by frame
    ret, frame = vid.read()

    # Display the resulting frame
    # putting the FPS count on the frame

    # time when we finish processing for this frame
    new_frame_time = time.time()
    # Calculating the fps

    # fps will be number of frame processed in given time frame
    # since their will be most of time error of 0.001 second
    # we will be subtracting it to get more accurate result
    fps = 1/(new_frame_time-prev_frame_time)
    prev_frame_time = new_frame_time

    # converting the fps into integer
    fps = int(fps)

    # converting the fps to string so that we can display it on frame
    # by using putText function
    fps = str(fps)
    cv2.putText(frame, "FPS:"+fps, (7, 70), font, 3, (100, 255, 0), 3,
    cv2.LINE_AA)
    cv2.imshow('frame', frame)

    # # Convert to grayscale
    gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY) #converts BGR to
gray

    # # Apply sobel filter in x and y directions
    grad_x = cv2.Sobel(gray, cv2.CV_64F, 1, 0, ksize=5) #x gradient
    grad_y = cv2.Sobel(gray, cv2.CV_64F, 0, 1, ksize=5) #y gradient

```

```

67
68
69 # # scale the gradient to -1 to 1 range
70 grad_x = grad_x / np.absolute(grad_x).max()
71 grad_x = np.uint8(128+127*grad_x)
72 grad_y = grad_y / np.absolute(grad_y).max()
73 grad_y = np.uint8(128+127*grad_y)
74
75
76 # computing the gradient magnitude and orientation
77 magnitude = np.sqrt((grad_x ** 2) + (grad_y ** 2))
78 orientation = np.arctan2(grad_y, grad_x) * (180 / np.pi) % 180 # converting to degree
79
80 #converting the arrays from float16 to uint8 to display
81 magnitude =magnitude.astype(np.uint8)
82 orientation =orientation.astype(np.uint8)
83
84 # time when we finish processing for this frame
85 new_frame_time = time.time()
86 # Calculating the fps
87
88 # fps will be number of frame processed in given time frame
89 # since their will be most of time error of 0.001 second
90 # we will be subtracting it to get more accurate result
91 fps = 1/(new_frame_time-prev_frame_time)
92 prev_frame_time = new_frame_time
93
94 # converting the fps into integer
95 fps = int(fps)
96
97 # converting the fps to string so that we can display it on frame
98 # by using putText function
99 fps = str(fps)
100 cv2.putText(magnitude, "FPS:"+fps, (7, 70), font, 3, (100, 255, 0),
101 3, cv2.LINE_AA)
102 cv2.imshow('Gradient magnitude',magnitude)
103
104 # Wait key for 1ms, if it is 'q' quit
105 if cv2.waitKey(1) & 0xFF == ord('q'):
106     break
107
108 # After the loop release the video capture object
109 vid.release()
110 # Destroy all the windows
111 cv2.destroyAllWindows()

```

Output:

